

Applying Python Architectural Design to The Scents Company

Problem Identification

Implementing the Architectural Patterns in Python (APP) book to practical application such as solving a domain problem close to one's own experience is the essence of this task. The refactoring approach opens flexibility when writing scripts, making room for expansion and improvement in the future without breaking the system. The idea of compartmentalizing functionalities into sections makes sense, especially when dealing with only a portion of the code, such as adding more features or replacing an existing application with something better. The principle of Test Driven Development (TDD) allows for a resilient code because it is essentially writing a script that has already gone through a series of tests. Efficiency also plays an important role in TDD, with tests already being written and available when the code is complete. The Scents Company, a fictitious company aiming to be the main producer of scents incorporated in filmmaking and videogaming entertainment industry, has a domain that requires an inventory management of its products similar to the App book's allocation system.

The entertainment industry, specifically filmmakers and video game developers, have been tremendously successful in delivering highly entertaining and realistic media, but where they still struggle is to incorporate the sense of smell with their movies and games for their audience to enjoy a fully immersive experience. No doubt that Hollywood and game creators have tried over the years, but to this day scent is still not mainstream in movies and games.

The idea has been around for ages, and there were some serious attempts that unfortunately failed along the way. The machines were complicated and expensive powered by compressed air containing various scents that were manually activated in certain movie scenes. The system was found to be rather noisy. Some areas of the theater were more saturated than others, and some viewers complained that the odors released were delayed. Clearly the technology needed was not available at that time.

Moviemakers have successfully satisfied the audiences' craving for realistic visual effects, as well as astounding sounds. Lately, innovators are incorporating body suits to deliver realistic sense of touch to the experience. This domain is clearly lucrative to filmmakers and video game makers, and the demand is definitely not going away.

Innovators incorporate scents into their own products in small scale. What would help is to develop a software system that could easily be programmed to calculate complex chemical mixtures, produce the odors cheaply, and be the major producer of scents for the entire entertainment industry. With economies-of-scale in mind, odors can be produced inexpensively, and with ease of use, scents in movies and video games can be mainstream. Another required part of the system is a software that runs and analyzes the logistics of the business demand. By running analytics, the software should be able to identify the most common scents applicable to movies and games, mass-produce them, and offer them cheaply to game creators and moviemakers. Also, startup innovators use their virtual reality goggles to dispense odor, so flexible scent dispensing designs must be kept in mind to offer to such customers. For movie theaters, scents can be released from the audiences' seats. Movie theaters are not required to

undergo major infrastructure upgrade to make this work. Each theater seat can be attached with scent-releasing device connected to vials of odor-producing chemicals, and the entire mechanism is controlled remotely via Bluetooth. A software can manage the release of odors timely for the entire theater as the movie is viewed. As audiences become more particular and critical, the software system must be easily updated to meet the changing demand. For instance, the sensation of ocean breeze scent varies per person, but running analytics from surveys collected can zero in on the most accepted scent of the ocean breeze and update the scent accordingly.

The scents are produced in the lab but have not been introduced for market research. There are expensive suppliers and cheap suppliers. The expensive supplier uses premium oils, while the cheap supplier uses lower grade oils, but both suppliers can produce the same list of scents (e.g., sea breeze, night club, jungle). The audience does not know which scents are produced by either supplier. The audience decides which scents are accepted and believable. The idea is to produce highly accepted scents as cheap as possible. The market research comes back with percentages of favorability per type of scent. The parameters mainly focus on the cheap supplier. If a cheap scent is 40% favored or higher, go with the cheap supplier. Otherwise go with the expensive supplier.

Collected data would look something like this:

scents	expensive supplier	cheap supplier	expensive costs	cheap costs
sea breeze	40%	60%	\$ 25.00	\$ 15.00
nightclub	35%	65%	\$ 18.00	\$ 16.00
jungle	70%	30%	\$ 30.00	\$ 25.00
forest	50%	50%	\$ 22.00	\$ 14.00
busy street	49%	51%	\$ 27.00	\$ 20.00

The business foundations mentioned above need to be fulfilled to start the Scents Company business, but once all these foundations have been established, the business will mainly focus on selling the scent products to customers, and this is when the principles from APP book can be implemented. Fulfilling orders based on inventory and shipment schedule is the focus hereafter, and the allocation system from the APP book will provide a system that manages the Scents Company's inventory and ordering system. Filling orders require proper inventory management, shipment control, and flawless timing to avoid out of stock mishaps, filling the wrong products, and shipment timing issues. Scents Company requires an allocation system that produces a list of its products, a batch system that manages its products, and a reliable ordering system for its customers. All these data should be stored in a database. The system must manage all these areas at all times. There should be a mechanism to allocate product batches from the inventory to fill customer orders. A quantity change functionality should allow for updating batch quantity, should customers change their order logistics. The system should be able to handle product recordkeeping such as using stock-keeping unit (SKU), quantity, and estimated time of arrival (ETA). The system should also be able to handle sending notifications both internally for employees and externally for customers.

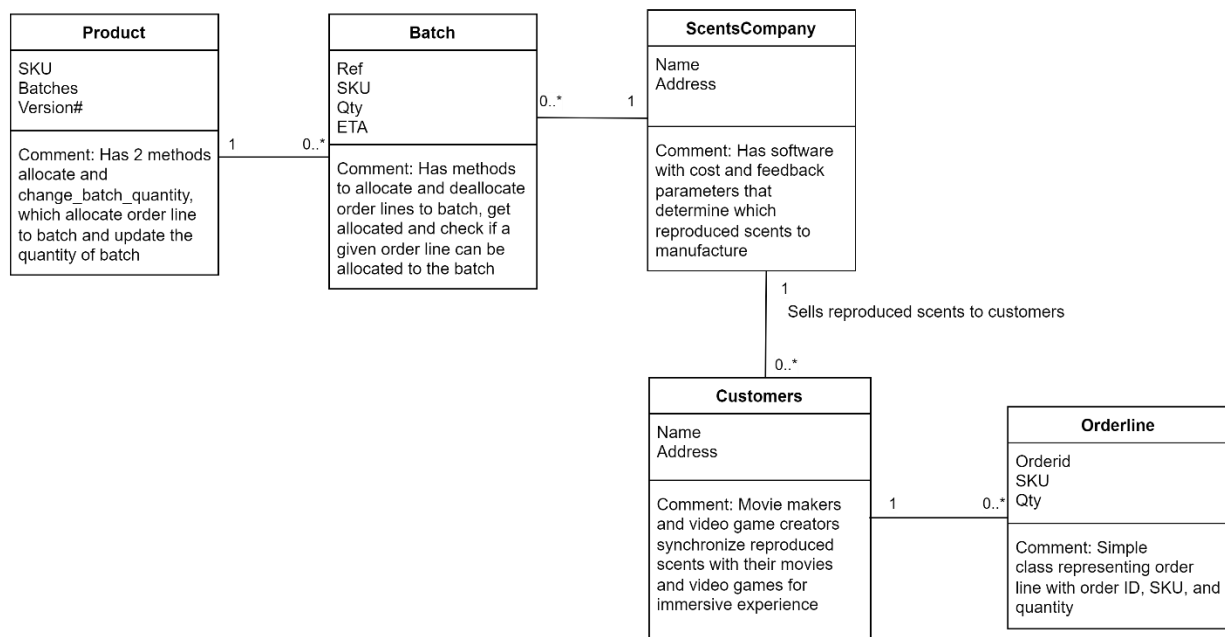
The APP book also emphasizes the importance of flexibility in the system to make room for improvements in the future. Encapsulation of domain logic is essential to encapsulate the domain logic and keep it separate and organized in a single module. Flexibility and extensibility allow for new commands and events to be added to the system without modifying the existing code. Testability allows

for testing scripts in isolation, making it easier to write unit tests for the domain model. Separation of concerns helps to make the code more modular and easier to maintain.

The use of a message bus improves modularity, scalability, and flexibility of the application for the Scents Company. It provides a centralized location to handle messages and decouples the handling of messages from the rest of the application's logic, which can simplify the code and make it more maintainable. This makes it easier to test and debug the message handling code. The use of a message bus also enables the implementation of complex domain logic that involves coordinating multiple events and commands, which can be difficult to do in a purely procedural application. It provides a way to handle events and commands asynchronously by queuing them up and processing them one at a time. It provides a way to handle any new events that are generated during the handling of events or commands.

Domain Modeling

Scents Company Revised Domain Model



Use Modeling

To apply the APP architecture in the Scents Company's allocation system, the company's inventory is set up as the Batch class, the list of products is the Product class, and the customer orders is the dataclass Orderline. The Product class represents a product and contains a list of batches of that product, a version number, and a list of events that have happened to the product. It has two methods, namely allocate and change_batch_quantity, which allocate an order line to a batch and update the quantity of a batch, respectively. The Batch class represents a batch of a product and contains a reference, SKU, quantity, ETA, a set of allocations, and methods to allocate and deallocate order lines to the batch. It also has methods to get allocated and available quantity and to check if a given order line

can be allocated to the batch. The OrderLine dataclass is a simple class that represents an order line with an order ID, SKU, and quantity.

The APP book's architecture provides event and command handlers for the Scents Company's allocation system which creates batches of products, allocates orders to available batches, and changes the quantity of a batch. The `add_batch` function creates a new batch or adds a batch to an existing product, the `allocate` function allocates an order to a batch, the `reallocate` function reallocates an order that was previously deallocated, and the `change_batch_quantity` function changes the quantity of a batch. The `send_out_of_stock_notification` function sends a notification when a product goes out of stock, the `publish_allocated_event` function publishes an allocated event to a message broker, and the `add_allocation_to_read_model` and `remove_allocation_from_read_model` functions add or remove an allocation from the read model. The script also defines two dictionaries, `EVENT_HANDLERS` and `COMMAND_HANDLERS`, which map event and command types to their respective handlers. When an event or command is received, the appropriate handler function is called to handle it.

Functional Test Plan

Unit Tests

Test Batches

This test contains a series of unit tests that handles order allocation. It tests for various methods of the Batch class in `allocation.domain.model` module and its ability to allocate and deallocate order lines.

The Batch class represents a group of items with a particular SKU that share an ETA. The class has methods to allocate and deallocate order lines, and to check if a given order line can be allocated to the batch based on its available quantity.

The unit tests check if the following functionalities of the Batch class work as expected. Allocating an order line reduces the available quantity of the batch by the quantity of the order line. The `can_allocate` logic returns True if the available quantity of the batch is greater than or equal to the required quantity of the order line, and False if not. The `can_allocate` logic returns False if the SKU of the order line does not match the SKU of the batch. Allocation of the same order line multiple times has no effect on the available quantity of the batch.

The first test called `test_allocating_to_a_batch_reduces_the_available_quantity` tests whether allocating a quantity to a batch reduces the available quantity of that batch by the same amount.

The next three tests labeled `test_can_allocate_if_available_greater_than_required`, `test_cannot_allocate_if_available_smaller_than_required`, and `test_can_allocate_if_available_equal_to_required` test whether a batch can allocate a line of a certain quantity based on its available quantity.

The fourth test named `test_cannot_allocate_if_skus_do_not_match` tests whether a batch can allocate a line only if the SKU of the batch matches that of the line.

The fifth test called `test_allocation_is_idempotent` tests whether allocating the same line to a batch twice will only reduce the available quantity of that batch once.

Overall, the tests ensure that the Batch class can properly allocate and deallocate order lines based on its available quantity and SKU.

Test Handlers

This test contains a set of unit tests that manages product allocations. The module handles commands related to product batches such as creating a batch, allocating from a batch, and changing the quantity of a batch. The unit tests ensure that the module works correctly in different scenarios.

The module defines a fake implementation of a repository and a unit of work, which are used to isolate the tests from the actual persistence mechanism. The module also defines a fake implementation of a notification service that records the notifications sent during the tests. The FakeRepository class provides an in-memory storage solution for products and their associated batches (i.e., groups of inventory units with the same SKU and creation date). The FakeUnitOfWork class provides an in-memory implementation of the unit of work pattern for coordinating changes to the repository. It includes a committed flag to indicate whether the changes have been persisted. The FakeNotifications class provides a simple notification service that records messages sent to each destination in a dictionary.

The module defines three classes, namely TestAddBatch, TestAllocate, and TestChangeBatchQuantity. Each of these classes contains one or more test methods that test the functionality of the corresponding commands. For example, TestAddBatch has two test methods test_for_new_product and test_for_existing_product. These methods create batches using the CreateBatch command and verify that the batch was created correctly. TestAllocate tests the Allocate command for reserving inventory units for a new order and updating the available quantity of the associated batch. TestChangeBatchQuantity tests the ChangeBatchQuantity command for adjusting the available quantity of a batch and reallocating inventory units as necessary.

The test methods use the bootstrap_test_app function to obtain an instance of the module's business logic. The function initializes the necessary dependencies (such as the fake repository and unit of work) and returns an instance of the business logic. A bootstrap_test_app function creates a test version of the application by calling the bootstrap.bootstrap function with the appropriate parameters (including the fake repository, unit of work, and notification classes).

The tests use the pytest framework to assert that the application behaves as expected. For example, the TestAllocate.test_allocates method creates a new batch, allocates some inventory units for an order, and verifies that the available quantity of the batch has been reduced accordingly. It uses the assert statement to check that the available_quantity attribute of the batch matches the expected value.

Overall, the module provides unit tests that test the functionality of the product allocation module. The tests ensure that the module works correctly in various scenarios and that it correctly handles errors and sends notifications when necessary.

Test Product

This module creates instances of the Product, Batch, and OrderLine classes and checks the behavior of their methods in different scenarios. Each test method creates an instance of Product and a corresponding OrderLine object, and then calls the allocate method on the Product instance, which attempts to allocate inventory from the product's batches to the order line. The tests assert various

properties of the Product and its batches before and after the allocation, as well as the creation of certain events related to the allocation process.

The `test_prefers_warehouse_batches_to_shipments` checks that the allocation system assigns an order line to an in-stock batch before considering a shipment batch. It creates a product with two batches of the same SKU, one with an ETA of None (representing in-stock inventory) and another with a future ETA (representing inventory that will be received from a supplier). It allocates an order line to the product and asserts that the allocation comes from the in-stock batch, leaving the shipment batch untouched.

The `test_prefers_earlier_batches` checks that the allocation system assigns an order line to the earliest batch that can fulfill it. It creates a product with three batches of the same SKU, each with a different ETA. It allocates an order line to the product and asserts that the allocation comes from the batch with the earliest ETA, leaving the other two batches untouched.

The `test_returns_allocated_batch_ref` checks that the `allocate` method of the Product class returns the reference of the batch that was allocated to the order line. It creates a product with two batches of the same SKU, one with an ETA of None and another with a future ETA. It allocates an order line to the product and asserts that the return value of the `allocate` method is the reference of the in-stock batch.

The `test_outputs_allocated_event` checks that an allocated event is produced when a batch is allocated to an order line. It creates a product with one batch of a particular SKU and allocates an order line to the product. It asserts that the last event in the product's event list is an allocated event with the correct properties.

The `test_records_out_of_stock_event_if_cannot_allocate` checks that an `OutOfStock` event is produced when there is no batch that can fulfill an order line. It creates a product with one batch of a particular SKU and allocates an order line to the product to exhaust the batch's available quantity. It attempts to allocate another order line for the same SKU to the product and asserts that the return value is None (since the allocation failed) and that the last event in the product's event list is an `OutOfStock` event with the correct SKU.

The `test_increments_version_number` checks that the version number of the Product instance is incremented when an order line is successfully allocated to a batch.`test_increments_version_number`. It creates a product with one batch of a particular SKU and sets its version number to 7. It allocates an order line to the product and asserts that the product's version number is incremented to 8.

Integration Tests

Test Repository

This module defines an abstract repository interface, `AbstractRepository`, and a concrete implementation of the repository using the SQLAlchemy Object-relational Mapper (ORM), `SqlAlchemyRepository`. The abstract repository interface declares methods for adding a product to the repository, getting a product by SKU, and getting a product by batch reference. The implementation of these methods is left to concrete subclasses of `AbstractRepository`. The implementation of `_add` adds the given product to the repository session. The implementation of `_get` queries the session for a Product

object with the given SKU. The implementation of `_get_by_batchref` queries the session for a Product object that has a batch with the given batch reference.

This module tests whether the `get_by_batchref` method of the `SqlAlchemyRepository` class in the `allocation.adapters.repository` module can retrieve the correct Product object based on the reference of one of its batches. The `SqlAlchemyRepository` implementation provides concrete implementations of the `_add`, `_get`, and `_get_by_batchref` methods using `SQLAlchemy` queries. First, it creates a few Batch and Product objects and adds them to the `SqlAlchemyRepository` instance. Then it asserts that calling `get_by_batchref` with the reference of one of the batches returns the corresponding Product object, and that calling it with the reference of another batch returns a different Product object. The `SqlAlchemyRepository` implementation also maintains a set of products that have been added to or queried from the repository.

Test Unit of Work

This module defines the following four test cases, each with a unique purpose, and a set of helper functions to facilitate the tests.

The `test_uow_can_retrieve_a_batch_and_allocate_to_it` checks whether a batch of products can be retrieved from the database, and whether an order line can be allocated to the batch. It does this by creating a new batch in the database, retrieving it using a unit of work (uow), allocating an order line to the batch, and then verifying that the correct batch was allocated to the order line.

The `test_rolls_back_uncommitted_work_by_default` checks whether uncommitted changes to the database are rolled back when a uow is exited without committing. It does this by creating a new batch in the database using a uow, and then verifying that the batch was not actually created in the database.

The `test_rolls_back_on_error` checks whether uncommitted changes to the database are rolled back when an error occurs during a uow. It does this by creating a new batch in the database using a uow and then raising an exception. It then verifies that the batch was not actually created in the database.

This module also defines the following several helper functions to facilitate the tests. The `insert_batch` inserts a new batch of products into the database. The `get_allocated_batch_ref` retrieves the batch reference for an order line that has already been allocated to a batch. The `try_to_allocate` attempts to allocate an order line to a batch using a new uow. The script also defines several constants, such as `random_sku`, `random_batchref`, and `random_orderid`, which are used to generate random values for testing purposes.

Test Views

This module contains the following two test functions that test the allocation functionality.

The `test_allocations_view` tests the `views.allocations` function. It first sets up a bus object using the `bootstrap` function from the `allocation` module and a `SQLite` session factory. The bus object is used to execute a series of commands that create batches of items and allocate orders for those items. After executing the commands, the test function calls the `views.allocations` function with an order ID and the unit of work object associated with the bus. The function returns a list of dictionaries containing

information about the allocated batches for the given order. The test function then asserts that the returned list is equal to an expected list of dictionaries containing information about the allocated batches for the given order.

The test_deallocation tests the deallocation of batches. It follows a similar setup as the first test function, but additionally changes the quantity of one of the created batches using the commands.ChangeBatchQuantity command. The test function then calls the views.allocations function again with the same order ID and unit of work object, and asserts that the returned list of allocated batches is equal to an expected list. In this case, the test expects only one batch to be allocated since the quantity of the other batch was reduced.

The script also defines a fixture called sqlite_bus which sets up a bus object for use in the test functions. The yield keyword is used to return the bus object to the test function and clear the mappers after the test function completes.

End-to-End Tests

Test API

As the Scents Company matches product batches to orders and returns appropriate error messages when invalid SKUs are requested, it is necessary to test that the API client functions as expected. The first test test_happy_path_returns_202_and_batch_is_allocated checks the successful allocation of a batch of products to an order. The test uses the api_client module to generate a random order ID and two random product SKUs. It creates three random batch references for the two SKUs. It adds the two product SKUs to the database along with their respective batches. It makes a request to allocate a quantity of the first product SKU to the order ID. It checks that the request returns a status code of 202 (Accepted). Lastly, it checks that the correct batch has been allocated to the order ID.

The second test function test_unhappy_path_returns_400_and_error_message checks that an error message is returned when an invalid SKU is requested. This test also uses the api_client module to generate a random order ID and an unknown product SKU. It makes a request to allocate a quantity of the unknown product SKU to the order ID, with the expect_success flag set to False to ensure that the request fails. It checks that the request returns a status code of 400 (Bad Request). It checks that the error message in the response body matches the expected error message. It checks that a subsequent request to retrieve the allocation for the order ID returns a status code of 404 (Not Found).

Overall, these tests ensure that the API client functions as expected, correctly allocating product batches to orders and returning appropriate error messages when invalid SKUs are requested.

Test External Events

The test_change_batch_quantity_leading_to_reallocation tests the scenario where a change in the quantity of an allocated batch leads to the reallocation of the order to a different batch. The test case begins by generating a random order ID and SKU, and two random batch references named earlier_batch and later_batch. The API client is then used to add two batches to the system with the specified references, SKUs, and quantities, and estimated times of arrival. The API client is then used to allocate units of the SKU to the order with the generated order ID. The allocation is verified by checking the batch reference of the allocated batch using the API client.

Testable Implementation

Adapters

Notifications

This module defines two classes that handle notifications in the allocation system. The first class is an abstract base class named `AbstractNotifications` that defines a single abstract method `send(destination, message)`. The `Abstract Base Class (abc)` module is used to declare this class as an abstract base class. The second class is named `EmailNotifications` and it extends the `AbstractNotifications` class. It provides an implementation for the `send` method that sends an email notification to the specified destination with the given message. The constructor of `EmailNotifications` takes two optional arguments, namely `smtp_host` and `port`, which default to the values obtained from the `config` module. This module is imported at the beginning of the script, and provides configuration settings for the allocation system, including email host and port. Inside the constructor, an SMTP server object is created using the provided `smtp_host` and `port` values. The `noop` method is called on the server object to check if the connection to the SMTP server is valid. The `send` method of `EmailNotifications` creates a message with a subject of "allocation service notification" and the specified message as the body. The message is then sent using the SMTP server object's `sendmail` method, with the `from_addr` and the `to_addrs` set to the specified destination. Finally, the first line of the script includes a pylint comment to disable the "too-few-public-methods" warning for the module. This warning is raised by pylint when a class has too few public methods. Disabling it in this case indicates that it is intentional and not a problem.

ORM

This module defines database tables and mappers for SQLAlchemy ORM. It starts with the necessary imports, creates a metadata object to store table and mapper information, and sets up logging. There are four tables defined in this script, namely `order_lines`, `products`, `batches`, and `allocations`. The `order_lines` table stores information about individual order lines, including the sku, quantity, and order ID. The `products` table stores information about individual products, including the sku and version number. The `batches` table stores information about batches of products, including a unique reference, the product sku, the purchased quantity, and an estimated time of arrival. The `allocations` table is a many-to-many join table that associates order lines with batches. There is also an `allocations_view` table defined, which is not backed by a database table but instead is a view that combines data from the `orders`, `batches`, and `allocations` tables to display allocation information. The `start_mappers` function is designed to set up the mappings between the domain models and the database tables. It defines three mappers: one for the `OrderLine` model, one for the `Batch` model, and one for the `Product` model. The `Product` mapper includes a relationship to the `Batch` mapper and the `Batch` mapper includes a relationship to the `OrderLine` mapper via the `_allocations` attribute. Finally, there is an event listener that is triggered whenever a `Product` is loaded from the database. It sets the `product's events` attribute to an empty list, which is used by the domain model to track changes to the `Product` object.

Repository

This script contains the implementation of the `AbstractRepository` class, an abstract interface for repositories that manage the persistence of domain objects. It also includes a concrete implementation of this interface using SQLAlchemy. The `AbstractRepository` defines three methods, namely `add`, `get`, and `_add`. The `add` method takes a `Product` object and adds it to the repository, while the `get` method retrieves a `Product` object from the repository given its SKU. The `_add` method is an abstract method that needs to be implemented by concrete repositories to define how a `Product` object is actually added to the repository.

The `SqlAlchemyRepository` class is a concrete implementation of the `AbstractRepository` interface that uses SQLAlchemy to persist `Product` objects. It takes a session object in its constructor, which is an instance of `sqlalchemy.orm.session.Session`. The `_add`, `_get`, and `_get_by_batchref` methods are implemented using SQLAlchemy queries to add, retrieve, and search for products respectively.

This attribute is defined in the `AbstractRepository` and is used to keep track of products that have been accessed by the repository during its lifetime. This is useful for implementing caches or other performance optimizations.

Overall, the purpose of this script is to define an abstract repository interface and a concrete implementation of this interface that uses SQLAlchemy to persist `Product` objects. This allows the application to interact with a persistence layer without needing to know the specific implementation details.

Domain

Commands

This module defines three data classes that represent commands that can be executed in a system. The `@dataclass` decorator is used to automatically generate a constructor and other methods for each class based on their attributes. The `Command` class is an abstract base class that is used as a parent class for the other classes. It does not have any attributes or methods of its own and is used only to identify that the other classes are commands. The first command class `Allocate` has three attributes, namely `orderid`, which is a string representing the ID of the order being allocated, `sku`, which is a string representing the SKU of the item being allocated, and `qty`, which is an integer representing the quantity of the item being allocated. The second command class is `CreateBatch` and has four attributes, namely `ref`, which is a string representing the reference ID of the batch being created, `sku`, which is a string representing the SKU of the item being batched, `qty`, which is an integer representing the quantity of the item being batched, and `ETA`, which is an optional attribute representing the estimated time of arrival for the batch. The `Optional` type hint is used to indicate that this attribute is not required. The third command class `ChangeBatchQuantity` has two attributes, namely `ref`, which is a string representing the reference ID of the batch being modified, and `qty`, which is an integer representing the new quantity of the batch. The script also includes a pylint directive to disable a specific warning about having too few public methods. This warning is likely related to the fact that the command classes do not have any methods other than the automatically generated ones from the `dataclass` decorator.

Events

This module defines three data classes that represent events that can occur in a system. The `@dataclass` decorator is used to automatically generate a constructor and other methods for each class based on their attributes. The `Event` class is an abstract base class that is used as a parent class for the other classes. It does not have any attributes or methods of its own and is used only to identify that the other classes are events. The first event `Allocated` has four attributes, namely `orderid`, which is a string representing the ID of the order that was allocated, `sku`, which is a string representing the SKU of the item that was allocated, `qty`, which is an integer representing the quantity of the item that was allocated, and `batchref`, which is a string representing the reference ID of the batch that the item was allocated from. The second event class `Deallocated` has three attributes, namely `orderid`, which is a string representing the ID of the order that was deallocated, `sku`, which is a string representing the SKU of the item that was deallocated, and `qty`, which is an integer representing the quantity of the item that was deallocated. The third event class `OutOfStock` has one attribute `sku`, which is a string representing the SKU of the item that is out of stock. The script also includes a `pylint` directive to disable a specific warning about having too few public methods. This warning is likely related to the fact that the event classes do not have any methods other than the automatically generated ones from the `dataclass` decorator.

Model

This module defines two classes, `Product` and `Batch`, as well as a `dataclass` `OrderLine` and imports a few other classes from the `Commands` and `Events` modules. The `Product` class represents a product and contains a list of batches of that product, as well as a version number and a list of events that have happened to the product. The `Product` class has two methods, namely `allocate` and `change_batch_quantity`. The `allocate` method takes an `OrderLine` and tries to allocate it to a batch of the product. If successful, it returns the reference of the batch that was allocated to. If there are no batches that can fulfill the order, an `OutOfStock` event is added to the list of events for the product. The `change_batch_quantity` method takes a batch reference and a new quantity and updates the quantity of the specified batch. If the new quantity causes the available quantity of the batch to become negative, the method tries to deallocate orders from the batch until the available quantity becomes non-negative. The `Batch` class represents a batch of a product and contains a reference, a SKU, a quantity, an ETA, a set of allocations, and methods to allocate and deallocate order lines to the batch. The `Batch` class also defines methods to get the allocated quantity and available quantity of the batch, as well as a method to check if a given `OrderLine` can be allocated to the batch. The `OrderLine` `dataclass` is a simple class that represents an order line with an order ID, a SKU, and a quantity. The `__future__` import enables the use of annotations in function signatures. The `unsafe_hash=True` parameter in the `@dataclass` decorator for `OrderLine` allows instances of the class to be hashable.

Entrypoints

Flask App

This module sets up a Flask web application with three endpoints for managing inventory allocations. The first endpoint namely `/add_batch`, is a POST request that creates a new batch of inventory. The endpoint expects a JSON payload with the following properties, namely `ref`, which is a reference string for the batch, `sku`, the SKU code for the product being added, `qty`, the quantity of the

product in the batch, and ETA (optional), an ISO-formatted date string representing the estimated time of arrival for the batch. The endpoint creates a `commands.CreateBatch` object from the JSON payload, and passes it to the `bus.handle` method for processing. If the command is successful, the endpoint returns an HTTP status code of 201 ("Created") with the message "OK". The second endpoint, `/allocate`, is a POST request that allocates inventory to an order. The endpoint expects a JSON payload with the following properties, namely `orderid`, a unique identifier for the order being allocated to, `sku`, the SKU code for the product being allocated, and `qty`, the quantity of the product being allocated. The endpoint creates a `commands.Allocate` object from the JSON payload, and passes it to the `bus.handle` method for processing. If the command is successful, the endpoint returns an HTTP status code of 202 ("Accepted") with the message "OK". If the SKU is invalid, the endpoint returns an HTTP status code of 400 ("Bad Request") with an error message. The third endpoint `/allocations/<orderid>` is a GET request that retrieves a list of allocations for a given order. The endpoint expects an `orderid` parameter in the URL. The endpoint calls the `views.allocations` function to retrieve the list of allocations for the given order from the database. If the order is not found, the endpoint returns an HTTP status code of 404 ("Not Found") with the message "not found". If the order is found, the endpoint returns an HTTP status code of 200 ("OK") with a JSON payload containing the allocations.

Service Layer

Handlers

This module defines the event and command handlers for the allocation system. The system receives commands to create batches of products, allocate orders to available batches, and change the quantity of a batch. The event handlers handle events such as when an order is allocated to a batch or when a product goes out of stock. The `add_batch` function creates a new batch for a product or adds a batch to an existing product. It takes a `CreateBatch` command and a unit of work object as arguments. The function retrieves the product associated with the SKU of the command and adds the new batch to its list of batches. If the product does not exist, it creates a new one with the given SKU and adds the batch to it. The `allocate` function allocates an order to a batch of products. It takes an `Allocate` command and a unit of work object as arguments. The function creates a new `OrderLine` object from the command and retrieves the product associated with the SKU of the line. If the product does not exist, it raises an `InvalidSku` exception. Otherwise, it calls the `allocate` method of the product, which attempts to allocate the order to one of its batches. The `reallocate` function reallocates an order that was previously deallocated. It takes a `Deallocated` event and a unit of work object as arguments. The function creates a new `Allocate` command from the event and calls the `allocate` function to allocate the order again. The `change_batch_quantity` function changes the quantity of a batch. It takes a `ChangeBatchQuantity` command and a unit of work object as arguments. The function retrieves the product associated with the batch reference of the command and calls its `change_batch_quantity` method to update the quantity of the batch. The `send_out_of_stock_notification` function sends a notification when a product goes out of stock. It takes an `OutOfStock` event and a notification object as arguments. The function calls the `send` method of the notification object to send an email to the stock management team. The `publish_allocated_event` function publishes an `Allocated` event to a message broker. It takes an `Allocated` event and a publish function as arguments. The function calls the `publish` function with the event type and the event data. The `add_allocation_to_read_model` function adds an allocation to the read model. It takes an `Allocated` event and a SQL Alchemy unit of work object as arguments. The function inserts the allocation data into the `allocations_view` table. The `remove_allocation_from_read_model` function

removes an allocation from the read model. It takes a Deallocated event and a SQL Alchemy unit of work object as arguments. The function deletes the allocation data from the allocations_view table. The script also defines two dictionaries, namely EVENT_HANDLERS and COMMAND_HANDLERS. These dictionaries map event and command types to their respective handlers. When an event or command is received, the appropriate handler function is called to handle it.

MessageBus

This module defines a MessageBus class that provides a way to handle messages in a message-driven architecture. The MessageBus class takes in an instance of a unit_of_work.AbstractUnitOfWork, a dictionary of event handlers, and a dictionary of command handlers. The handle method of the MessageBus class takes a Message object, which can either be a commands.Command or an events.Event. The MessageBus handles the message by first adding it to a queue and then processing the messages in the queue one by one. For each message, the handle method determines whether it is a command or an event and calls either the handle_command or handle_event method, respectively. The handle_event method looks up the appropriate handler function for the given event type in the event_handlers dictionary and calls it, passing in the event as an argument. If there are multiple handlers registered for the event, all of them are called in the order they appear in the list. If an exception is raised during the handling of the event, it is caught, logged, and the loop continues to the next event handler. The handle_command method looks up the appropriate handler function for the given command type in the command_handlers dictionary and calls it, passing in the command as an argument. If an exception is raised during the handling of the command, it is caught, logged, and re-raised. Both the handle_command and handle_event methods collect any new events generated during their respective processing and add them to the message queue.

Overall, the MessageBus class provides a way to loosely couple message producers and consumers in a message-driven architecture by abstracting the handling of messages away from the individual components of the system.

Unit of Work

This module manages a unit of work using the SQLAlchemy ORM library. The AbstractUnitOfWork abstract base class defines the interface for a unit of work, which provides a consistent view of the database to the application. It has two abstract methods, namely _commit and rollback. It also has an attribute products, which is of type repository.AbstractRepository. This is a placeholder attribute that is expected to be implemented by subclasses to provide access to the domain objects being persisted. The SQLAlchemyUnitOfWork is a concrete implementation of AbstractUnitOfWork. It takes a session_factory argument that defaults to DEFAULT_SESSION_FACTORY, which is a sessionmaker object that creates a new Session object that is bound to a database engine specified in the config module. The SQLAlchemyUnitOfWork overrides the __enter__ and __exit__ methods to create a new Session and to close the session, respectively. The __enter__ method also initializes the products attribute with a SQLAlchemyRepository object that is bound to the session. The _commit method commits the changes made in the session to the database, and the rollback method rolls back any uncommitted changes made in the session. The collect_new_events method of AbstractUnitOfWork is used to collect any new domain events that are created during the execution of the unit of work. This method yields any events that have not been processed yet from the events list of each Product object in the products attribute of the unit of work.

Bootstrap

This module defines a bootstrap function that initializes a message bus with event and command handlers and injects their dependencies. It also defines a helper function `inject_dependencies` to perform the injection. The bootstrap function takes several optional parameters. The `start_orm` indicates whether to start the ORM and is `True` by default. The `uow` is an instance of `unit_of_work.AbstractUnitOfWork`, which is `unit_of_work.SqlAlchemyUnitOfWork` by default. The `notifications` is an instance of `AbstractNotifications`, which is `EmailNotifications` by default. The function first checks if the `notifications` parameter is `None` and initializes it with an instance of `EmailNotifications` if it is. Then, if `start_orm` is `True`, it calls `orm.start_mappers` to start the ORM. The function then creates a dependencies dictionary with the `uow`, `notifications`, and `publish` parameters, which will be injected into the handlers. Next, it uses a dictionary comprehension to create a new dictionary of `injected_event_handlers`. The keys of this dictionary are the event types, and the values are lists of handlers with their dependencies injected. To do this, it iterates over the `EVENT_HANDLERS` dictionary in the `handlers` module, which maps event types to handler functions. For each event type and its handlers, it calls the `inject_dependencies` function on each handler, passing in the dependencies dictionary. The resulting list of injected handlers is then stored in the `injected_event_handlers` dictionary under the corresponding event type. Similarly, the function creates an `injected_command_handlers` dictionary that maps command types to injected handler functions, using a dictionary comprehension and the `inject_dependencies` function. Finally, the function creates and returns a `messagebus.MessageBus` instance with the `uow`, `injected_event_handlers`, and `injected_command_handlers` as parameters. The `inject_dependencies` function takes a handler function and a dependencies dictionary, and returns a new lambda function that calls the handler with the given dependencies injected as keyword arguments. It does this by inspecting the signature of the handler function, and finding the parameters that match the keys in the dependencies dictionary. It then creates a new dictionary of the matching parameters and their corresponding dependencies and passes this dictionary to the handler using keyword argument unpacking. The resulting lambda function is returned as the injected handler.

Config

This module provides functions to retrieve the necessary configuration parameters from environment variables for the allocation system. The `get_api_url` returns the API URL of the system. The function retrieves the `API_HOST` environment variable to get the hostname of the API server. If this environment variable is not set, the default value of `localhost` is used. The function sets the port to `5005` if the hostname is `localhost`, otherwise it sets the port to `80`. The function sets the port to `63791` if the hostname is `localhost`, otherwise it sets the port to `6379`. The `get_email_host_and_port` returns a dictionary with the email server host and port configuration. The function retrieves the `EMAIL_HOST` environment variable to get the hostname of the email server. If this environment variable is not set, the default value of `localhost` is used. The function sets the SMTP port to `11025` if the hostname is `localhost`, otherwise it sets the SMTP port to `1025`. The function sets the HTTP port to `18025` if the hostname is `localhost`, otherwise it sets the HTTP port to `8025`.

Views

This module defines a function `allocations` that takes an `orderid` string and a `SqlAlchemyUnitOfWork` object as parameters. It queries the database using the `uow` object to get the `sku` and `batchref` values for all allocations associated with the given `orderid`. The function returns a list of

dictionaries where each dictionary contains a sku and batchref key-value pair. The function is using a context manager (`with`) to automatically handle committing and rolling back the changes made during the database transaction. The uow object is a SQLAlchemyUnitOfWork which is a concrete implementation of an abstract class AbstractUnitOfWork that defines the contract for interacting with the database. The AbstractUnitOfWork class has a method `_commit` that is used to commit the changes made during the transaction and a method `rollback` to roll back the changes in case of any errors. The SQL query is a parameterized query that uses named placeholders to avoid SQL injection attacks. The placeholders are replaced with actual values at runtime using a dictionary that maps the placeholder names to their corresponding values. The query returns a result set that is converted to a list of dictionaries using a list comprehension. Each dictionary in the list represents an allocation and contains sku and batchref key-value pairs.

Overall, this module provides a way to retrieve allocation data from the database using a transactional approach with automatic rollback and commit features.

Setup

This module is a basic setup configuration file for the allocation package. It imports the setup function from `setuptools` and uses it to specify the package name, version, and packages included in the distribution. The name parameter specifies the name of the package, which in this case is `allocation`. The version parameter specifies the version number of the package, which in this case is `0.1`. The packages parameter specifies a list of packages to be included in the distribution. In this case, the only package included is the top-level allocation package.

Validation and Justification

The APP architecture is essential in implementing an efficient, methodical, and organized inventory management and ordering system for the Scents Company. The following enumerate the architectural benefits of the APP architecture for the Scents Company.

Domain

Entity and Aggregate Benefits

The Product and Batch are two entity components. Product represents a product that has a SKU and is composed of one or more Batch instances. It contains the logic to allocate a Batch to an OrderLine and to change the quantity of a Batch. It also keeps track of the events related to it, such as when it is allocated or out of stock. Batch represents a batch of products that share the same SKU and ETA. It contains the logic to allocate and deallocate OrderLine instances, and to determine if it can allocate an OrderLine. It also keeps track of its own state, such as the purchased and allocated quantity, and the ETA.

The Product and Batch components have several benefits, including:

1. **Modularity:** By encapsulating the logic related to products and batches into separate classes, the code becomes more modular, easier to read and maintain.
2. **Reusability:** The Product and Batch classes can be reused in other parts of the codebase, which reduces the amount of duplicated code and makes it easier to add new features or modify existing ones.

3. **Abstraction:** By abstracting away the details of the product and batch logic into separate classes, the rest of the codebase can interact with them through well-defined interfaces, which makes it easier to reason about the code and reduces the likelihood of bugs.
4. **Scalability:** The Product and Batch classes provide a solid foundation for scaling the system, as they allow for easy management of large numbers of products and batches and provide a clear separation of concerns.
5. **Testability:** The Product and Batch classes can be easily tested in isolation, which makes it easier to write comprehensive tests for the codebase and catch bugs early on in the development process.

Value Object Benefits

The OrderLine class in the Domain module is a value object because it is immutable and its equality is based on its attributes. This is indicated by the `@dataclass(unsafe_hash=True)` decorator, which automatically generates the `__eq__` and `__hash__` methods for the class. Additionally, the OrderLine instances are used as elements of the `set_allocations` in the Batch class, which further reinforces their immutability.

1. **Immutability:** As a value object, instances of OrderLine are immutable, which means that their state cannot be changed after they are created. This ensures that they remain consistent throughout the application and prevents accidental modification of their values.
2. **Type Hinting:** By defining the attributes of the OrderLine class, one can provide better type hinting for the code. This makes it easier for other developers to understand the expected data types and enables better type checking at compile time.
3. **Safety:** Because OrderLine is immutable, it is safer to use in multi-threaded or concurrent applications, where shared mutable objects can cause data race conditions and other concurrency issues.
4. **Clarity:** The OrderLine class provides a clear and concise way to represent an order line item, which makes the code more readable and easier to understand. It also enables one to encapsulate the logic related to order lines in a single class, which makes it easier to maintain and modify over time.

Events Benefits

This feature defines three events, namely Allocated, Deallocated, and OutOfStock. These events are used to model the state of Scents Company's inventory system. By using such events, Scents Company can keep track of changes to the inventory in a consistent and scalable way, which can help them make better decisions about how to allocate resources and fulfill orders. The Allocated event can be raised when a batch of products is reserved for an order, the Deallocated event can be raised when a batch is returned to the inventory, and the OutOfStock event can be raised when there is not enough inventory to fulfill an order. By defining events as dataclasses, it makes it easy to create new instances of the event with appropriate data, and to serialize or deserialize them as needed. Additionally, by inheriting from a base Event class, it makes it easier to identify and filter events from other types of messages or data that might be flowing through the system. Overall, this function provides a simple and straightforward way to define and work with events. This can make it easier to understand the data being passed between different parts of the system and can help prevent errors that can occur when passing data between different parts of the system in different formats.

Commands Benefits

By using this pattern, the implementation of an action can be decoupled from the request that triggers it, making it easier to modify, extend and maintain the code. By defining commands as separate classes, it allows for easier separation of concerns in the codebase. Commands represent the intention of the user or system and can be passed around between different parts of the application without needing to know the implementation details. This makes it easier to reason about the behavior of the system and maintain a consistent state. The use of dataclasses also simplifies the creation of instances of these command classes. It also allows for type annotations to be added to each field, which can help with static analysis and debugging. Also, using dataclasses to define commands can provide several benefits:

1. The code is more readable because the attributes of a command are clearly defined and documented in one place.
2. Dataclasses provide a concise way to define immutable objects, which can make the code safer and easier to reason about.
3. Dataclasses can be easily serialized to and deserialized from JSON, which can be useful in distributed systems.

Service Layer

Handler Benefits

This feature defines handlers for various commands and events in the allocation domain of the Scents Company. The main benefit of this script is that it decouples the implementation details of the domain from the rest of the application, making it easier to maintain and evolve over time. This feature benefits the Scents Company Overall because it encapsulates the business logic of the allocation domain in a modular and extensible way. It also allows for greater testability and scalability by separating concerns and minimizing dependencies between different parts of the system.

1. Encapsulation of domain logic: Each function in the script represents a specific operation that can be performed in the domain model, such as adding a new batch or allocating an order line. This helps to encapsulate the domain logic and keep it organized in a single module.
2. Flexibility and extensibility: The use of dictionaries (COMMAND_HANDLERS and EVENT_HANDLERS) allows for a flexible and extensible way to map commands and events to their corresponding handlers. This means that new commands and events can be added to the system without modifying the existing code, by simply adding a new handler function to the appropriate dictionary.
3. Testability: The functions in this script are designed to be testable in isolation, which can make it easier to write unit tests for the domain model. For example, the add_batch function takes a unit_of_work.AbstractUnitOfWork object as an argument, which can be replaced with a mock object during testing.
4. Separation of concerns: The functions in this script are responsible for handling specific types of commands and events and are not concerned with the implementation details of the system as a whole. This separation of concerns helps to make the code more modular and easier to maintain.

Unit of Work Benefits

The benefit of this feature is that it provides the Scents Company a standardized way to manage transactions and access the database within the allocation system using SQLAlchemy, which can help simplify the code and reduce the risk of errors when interacting with the database. This helps to ensure data integrity and consistency. The unit of work pattern can also simplify database code by abstracting away the details of the transaction management and providing a clean interface for working with the database. This function provides a useful abstraction for managing transactions and database access within the allocation domain, and makes it easier to write clean, testable code that interacts with a database.

Message Bus (Internal) Benefits

The benefits of using a message bus include improved modularity, scalability, and flexibility of the application for the Scents Company. It provides a centralized location to handle messages and decouples the handling of messages from the rest of the application's logic, which can simplify the code and make it more maintainable. This makes it easier to add new event and command handlers without modifying the application's core logic and makes it easier to test and debug the message handling code. Additionally, the use of a message bus enables the implementation of complex domain logic that involves coordinating multiple events and commands, which can be difficult to do in a purely procedural application. It provides a way to handle events and commands asynchronously by queuing them up and processing them one at a time. It provides a way to handle any new events that are generated during the handling of events or commands.

1. Decoupling: The MessageBus class allows for loose coupling between message producers and consumers, which can be helpful in large and complex systems. Producers only need to publish messages to the bus, without having to know who will handle them or how.
2. Scalability: A message-driven architecture using a MessageBus can scale horizontally, by adding more message consumers as needed.
3. Separation of concerns: By separating the handling of commands and events, the MessageBus class makes it easier to reason about the different types of messages and their handling. It also allows for different handlers to be used for different types of messages, based on the needs of the system.

Adapters

Repository Benefits

The benefit of having Repository for the Scents Company is that it provides an abstract repository class and a concrete implementation of the repository using SQL Alchemy. This implementation allows for adding and retrieving products by SKU and batch reference. It also tracks the products that have been added or retrieved so that it does not need to fetch them again from the database. This can improve performance and reduce unnecessary database calls. Using an abstract repository class allows for easy swapping of the underlying storage implementation, which is useful when switching from one database to another or when testing the code. It also makes the code more modular and easier to understand, as the implementation details of the database are hidden behind the abstract repository interface.

Event Publisher Benefits

This script defines a publish function that publishes events to a message broker. The benefit of using a message broker to the Scents Company is that it decouples the different parts of the system, allowing them to operate independently and asynchronously. This means that when an event occurs, it can be published to a message broker without waiting for any other part of the system to process it, improving the overall performance and responsiveness of the system. The Scents Company needs a publish function that is reliable, scalable, and flexible.

1. **Event-driven architecture:** The ability to publish events to a message broker is a key feature of an event-driven architecture. This allows for loose coupling between different parts of a system, where each component can subscribe to relevant events and respond accordingly.
2. **Scalability:** A message broker can be a high-performance, distributed data store that can handle large volumes of data and requests. By using a message broker, this script allows for horizontal scaling of the system, as multiple instances of the application can be deployed to handle different parts of the system.
3. **Flexibility:** The use of a message broker allows for flexibility in the system design, as new components can be added or removed without affecting the existing system.
4. **Reliability:** A message broker can be a robust and reliable data store that can handle high volumes of data and requests. This ensures that messages are delivered reliably to the subscribers, even in high-load scenarios.
5. **Logging:** The script uses the Python logging module to log messages at the INFO level when events are published. This provides visibility into the system and can be used for debugging or monitoring purposes.

Entrypoints

Web Benefits

Using the Flask web application that interacts with the allocation service, it allows the Scents Company user to add a new batch of products, allocate a specific quantity of a product to an order, and view allocations for a specific order. This function provides a convenient way for users to interact with the warehouse management system using a simple HTTP API.

Event Consumer Benefits

This feature implements a message broker pub/sub mechanism, which allows multiple processes to communicate with each other in a publish-subscribe pattern. The main benefit of this function is that it enables a distributed system to react to events or changes happening in other parts of the system in real-time, without the need for constant polling or waiting for responses. This feature allows for a scalable, asynchronous, and decoupled architecture, where different components of the system can communicate with each other through a message broker channels without directly knowing about each other.

Github repository link for the final project:

<https://github.com/asimbajon1/FinalProject-b.git>

Github link showing all work for the entire class:

<https://github.com/asimbajon1?tab=repositories>